

Module – 3

- **Operators:**

- Combining & splitting – UNION, SPLIT
- Sorting – ORDER BY, LIMIT
- Grouping Operator – GROUP, CO-GROUP
- Joining Operator - JOIN(INNER, SELF JOIN)

- **Pig Latin Built-in functions:**

- Eval functions (Avg, Max, Min, Sum, Count, Size, Concat, Tokenize)
- Bag & Tuple Functions
- String Functions
- Math Functions

- **Apache Pig - Running Scripts:**

- Creating pig script
- Commenting pig script
- Executing –running pig script – with/without parameters
- Sample examples

```
[cloudera@localhost ~]$ gedit dr1
```

```
Milan,1001,5,apollo,500  
Jay,1002,10,apollo,500  
lalit,1003,20,manipal,500  
Mohit,1004,15,columbia,600  
Chauhan,1005,30,narayana,550  
Suraj,1006,25,manipal,650
```

```
[cloudera@localhost ~]$ gedit dr2
```

```
meena,2001,20,rxdx,650  
leena,2001,15,st johns,450  
sonam,2002,30,rxdx,600
```

```
[cloudera@localhost ~]$ gedit empy1
```

```
7001,ameena,10,bang  
7002,amit,20,chennai  
7003,anand,30,bang  
7004,alen,15,hyd  
7005,alester,10,hyd  
7006,anshul,5,Chennai
```

```
[cloudera@localhost ~]$ gedit pnt1
```

```
101,harinath,5,domlur,1004  
102,nagarjun,10,varthur,1005  
103,chirajeevi,20,HAL,1006  
104,tarun,25,HSR,1004  
105,prabas,15,marthahalli,1006  
106,chaitanya,30,belandur,1003  
107,nani,27,krpuram,1004
```

```
[cloudera@localhost ~]$ pig -x local
```

```
grunt> clear
```

```
grunt> doc1 = load '/home/cloudera/dr1' using PigStorage(',') as
```

```
    (name:chararray, id:int, exp:int, hosp:chararray, fees:int);
```

```
grunt> dump doc1;
```

```
grunt> doc2 = load '/home/cloudera/dr2' using PigStorage(',') as
```

```
(name:chararray, id:int, exp:int, hosp:chararray, fees:int);
```

```
grunt> dump doc2;
```

UNION OPERATOR

The **UNION** operator of Pig Latin is used to merge the content of two relations. To perform UNION operation on two relations, their columns and domains must be identical.

Syntax

```
grunt> Relation_name3 = UNION Relation_name1, Relation_name2;
```

```
grunt> result = union doc1,doc2;
```

```
grunt> dump result;
```

```
(Milan,1001,5,apollo,500)
```

```
(Jay,1002,10,apollo,500)
```

```
(lalit,1003,20,manipal,500)
```

```
(Mohit,1004,15,columbia,600)
```

```
(Chauhan,1005,30,narayana,550)
```

```
(Suraj,1006,25,manipal,650)
```

```
(,,,) 
```

```
(meena,2001,20,rxdx,650)
```

```
(leena,2001,15,st johns,450)
```

```
(sonam,2002,30,rxdx,600)
```

SPLIT OPERATOR

The **SPLIT** operator is used to split a relation into two or more relations.

Syntax

Given below is the syntax of the **SPLIT** operator.

```
grunt> SPLIT Relation1_name INTO Relation2_name IF (condition1), Relation3_name IF(condition2)
```

```
grunt> split doc1 into senior if exp>15,junior if (exp>5 and exp<=15);
```

grunt> dump senior;

(lalit,1003,20,manipal,500)

(Chauhan,1005,30,narayana,550)

(Suraj,1006,25,manipal,650)

grunt> dump junior;

(Jay,1002,10,apollo,500)

(Mohit,1004,15,columbia,600)

ORDER BY OPERATOR

The **ORDER BY** operator is used to display the contents of a relation in a sorted order based on one or more fields.

Syntax

Given below is the syntax of the **ORDER BY** operator.

```
grunt> Relation_name2 = ORDER Relatin_name1 BY Field (ASC|DESC);
```

grunt> a = order doc1 by name asc;

grunt> dump a;

(Chauhan,1005,30,narayana,550)

(Jay,1002,10,apollo,500)

(Milan,1001,5,apollo,500)

(Mohit,1004,15,columbia,600)

(Suraj,1006,25,manipal,650)

(lalit,1003,20,manipal,500)

LIMIT OPERATOR

The **LIMIT** operator is used to get a limited number of tuples from a relation.

Syntax

```
grunt> Result = LIMIT Relation_name required number of tuples;
```

grunt> a = limit doc1 2;

grunt> dump a;

(Jay,1002,10,apollo,500)

(Milan,1001,5,apollo,500)

GROUP OPERATOR

The **GROUP** operator is used to group the data in a relation. It collects the data having the same key.

Syntax

```
grunt> Group_data = GROUP Relation_name BY age;
```

Q) Display the details of the doctors hospital wise.

```
grunt> gr = group doc1 by hosp;
```

```
grunt> dump gr;
```

```
(apollo,{(Milan,1001,5,apollo,500),(Jay,1002,10,apollo,500)})  
(manipal,{(lalit,1003,20,manipal,500),(Suraj,1006,25,manipal,650)})  
(columbia,{(Mohit,1004,15,columbia,600)})  
(narayana,{(Chauhan,1005,30,narayana,550)})
```

Grouping By Multiple Columns

Q) Display the details of the doctors hospital wise with same fees

```
grunt> a = group doc1 by (hosp,fees);
```

```
grunt> dump a;
```

```
((apollo,500),(Milan,1001,5,apollo,500),(Jay,1002,10,apollo,500))  
((manipal,500),(lalit,1003,20,manipal,500))  
((manipal,650),(Suraj,1006,25,manipal,650))  
((columbia,600),(Mohit,1004,15,columbia,600))  
((narayana,550),(Chauhan,1005,30,narayana,550))
```

CO-GROUP OPERATOR

The **COGROUP** operator works more or less in the same way as the GROUP operator. The only difference between the two operators is that the **group** operator is normally used with one relation, while the **cogroup** operator is used in statements involving two or more relations.

```
grunt> emp1 = load '/home/cloudera/empty1' using PigStorage(',') as  
          (id:int, name:chararray, exp:int, place:chararray);
```

```
grunt> a = cogroup doc1 by exp,emp1 by exp;
```

```
grunt> dump a;
```

```
(5,{(Milan,1001,5,apollo,500)},{(7006,anshul,5,chennai)})  
(10,{(Jay,1002,10,apollo,500)},{(7001,ameena,10,bang),(7005,alester,10,hyd)})
```

```
(15, {(Mohit,1004,15,columbia,600)}, {(7004,alen,15,hyd)})
(20, {(lalit,1003,20,manipal,500)}, {(7002,amit,20,chennai)})
(25, {(Suraj,1006,25,manipal,650)}, {})
(30, {(Chauhan,1005,30,narayana,550)}, {(7003,anand,30,bang)})
```

JOIN OPERATOR

The JOIN operator is used to combine records from two or more relations.

Self-Join

Self-join is used to join a table with itself as if the table were two relations.

```
grunt> doc = load '/home/cloudera/dr1' using PigStorage(',') as
      (name:chararray, id:int, exp:int, hosp:chararray, fees:int);
```

```
grunt> a = join doc by id, doc1 by id;
grunt> dump a;
(Milan,1001,5,apollo,500,Milan,1001,5,apollo,500)
(Jay,1002,10,apollo,500,Jay,1002,10,apollo,500)
(lalit,1003,20,manipal,500,lalit,1003,20,manipal,500)
(Mohit,1004,15,columbia,600,Mohit,1004,15,columbia,600)
(Chauhan,1005,30,narayana,550,Chauhan,1005,30,narayana,550)
(Suraj,1006,25,manipal,650,Suraj,1006,25,manipal,650)
```

Inner-Join

It is also referred to as equijoin. An inner join returns rows when there is a match in both tables.

```
grunt> pat = load '/home/cloudera/pnt1' using PigStorage(',') as
      (pid:int,name:chararray,age:int,addr:chararray,docid:int);
```

Q)display entire details of patient and their corresponding doctor

```
grunt> a = join doc1 by id, pat by docid;
grunt> dump a;
(lalit,1003,20,manipal,500,106,chaitanya,30,belandur,1003)
(Mohit,1004,15,columbia,600,101,harinath,5,domlur,1004)
(Mohit,1004,15,columbia,600,104,tarun,25,HSR,1004)
(Mohit,1004,15,columbia,600,107,nani,27,krpuram,1004)
(Chauhan,1005,30,narayana,550,102,nagarjun,10,varthur,1005)
(Suraj,1006,25,manipal,650,103,chirajeevi,20,HAL,1006)
(Suraj,1006,25,manipal,650,105,prabas,15,marthahalli,1006)
```

Pig Latin Built-in functions:

Eval functions (Avg, Max, Min, Sum, Count, Size, Concat, Tokenize)

```
[cloudera@localhost ~]$ gedit dr1
```

```
Milan,1001,5,apollo,500
Jay,1002,10,apollo,500
lalit,1003,20,manipal,500
Mohit,1004,15,columbia,600
Chauhan,1005,30,narayana,550
Suraj,1006,25,manipal,650
```

```
[cloudera@localhost ~]$ pig -x local
```

```
grunt> clear
```

```
grunt> doc = load '/home/cloudera/dr1' using PigStorage(',') as
```

```
(name:chararray, id:int, exp:int, hosp:chararray, fees:int);
```

```
grunt> dump doc;
```

```
(Milan,1001,5,apollo,500)
(Jay,1002,10,apollo,500)
(lalit,1003,20,manipal,500)
(Mohit,1004,15,columbia,600)
(Chauhan,1005,30,narayana,550)
(Suraj,1006,25,manipal,650)
```

Group All

You can group a relation by all the columns as shown below.

```
grunt> group_all = GROUP relation_name All;
```

```
grunt> gr = group doc all;
```

```
grunt> dump gr;
```

```
(all,{(Milan,1001,5,apollo,500),(Jay,1002,10,Apollo,500),(lalit,1003,20,manipal,500),(Mohit,15,1004,15,)
,(Chauhan,1005,30,narayana,550),(Suraj,1006,25,manipal,650),(Jay,102,10,apollo,50)})}
```

AVG():To compute the average of the numerical values within a bag.

Q)Display hospital name, fees and average fees among all the hospital.

```
grunt> result = foreach gr generate doc.hosp,doc.fees,AVG(doc.fees);  
({(apollo),(apollo),(manipal),(columbia),(narayana),(manipal),()},{(500),(500),(500),(600),(550),(650),()},{550.0})
```

MAX():To calculate the highest value for a column in a single-column bag.

Q)Display hospital name, fees and maximum fees among all the hospital.

```
grunt> result = foreach gr generate doc.hosp,doc.fees,MAX(doc.fees);  
grunt> dump result;  
({(apollo),(apollo),(manipal),(columbia),(narayana),(manipal),()},{(500),(500),(500),(600),(550),(650),()},{650})
```

MIN():To get the minimum (lowest) value (numeric or chararray) for a certain column in a single-column bag.

Q)Display hospital name, fees and minimum fees among all the hospital.

```
grunt> result = foreach gr generate doc.hosp,doc.fees,MIN(doc.fees);  
grunt> dump result;  
({(apollo),(apollo),(manipal),(columbia),(narayana),(manipal),()},{(500),(500),(500),(600),(550),(650),()},{500})
```

SUM():To get the total of the numeric values of a column in a single-column bag.

Q)Display hospital name, fees and total fees among all the hospital.

```
grunt> result = foreach gr generate doc.hosp,doc.fees,SUM(doc.fees);  
grunt> dump result;  
({(apollo),(apollo),(manipal),(columbia),(narayana),(manipal),()},{(500),(500),(500),(600),(550),(650),()},{3300})
```

COUNT(): To get the the number of tuples in a bag.

Q)Display total no:of tuples/rows in relation.

```
grunt> result = foreach gr generate COUNT(doc.id);  
grunt> dump result;  
(6)
```


SIZE(): To compute the number of elements based on any Pig data type.

Q)Display doctor name along with the length of doctor name in each row.

```
grunt> ans = foreach doc generate name,SIZE(name);
```

```
grunt> dump ans;
```

(Milan,5)

(Jay,3)

(lalit,5)

(Mohit,5)

(Chauhan,7)

(Suraj,5)

CONCAT(): To concatenate two or more expressions of same type.

```
grunt> ans = foreach doc generate CONCAT(name,hosp);
```

```
grunt> dump ans;
```

(Milanapollo)

(Jayapollo)

(lalitmanipal)

(Mohitcolumbia)

(Chauhannarayana)

(Surajmanipal)

Bag & Tuple Functions

TUPLE CONSTRUCTION:

```
grunt> a = foreach doc generate name,id,exp;
```

```
grunt> dump a;
```

(Milan,1001,5)

(Jay,1002,10)

(lalit,1003,20)

(Mohit,1004,15)

(Chauhan,1005,30)

(Suraj,1006,25)

BAG CONSTRUCTION:

```
grunt> a = foreach doc generate {(name,id,exp)},{name,id,exp};
```

```
grunt> dump a;
```

```
((Milan,1001,5)),((Milan),(1001),(5))  
((Jay,1002,10)),((Jay),(1002),(10))  
((lalit,1003,20)),((lalit),(1003),(20))  
((Mohit,1004,15)),((Mohit),(1004),(15))  
((Chauhan,1005,30)),((Chauhan),(1005),(30))  
((Suraj,1006,25)),((Suraj),(1006),(25))
```

MAP CONSTRUCTION:

```
grunt> a = foreach doc generate [name,exp];  
grunt> dump a;
```

```
([Milan#5])  
([Jay#10])  
([lalit#20])  
([Mohit#15])  
([Chauhan#30])  
([Suraj#25])
```

STRING BUILT IN FUNCTIONS

SUBSTRING()

Returns a substring from a given string.

Syntax:

SUBSTRING(string, startIndex, ending index+1)

```
grunt> ans = foreach doc generate (id,name),SUBSTRING (name, 0 , 2);  
grunt> dump ans;  
((1001,Milan),Mi)  
((1002,Jay),Ja)  
((1003,lalit),la)  
((1004,Mohit),Mo)  
((1005,Chauhan),Ch)  
((1006,Suraj),Su)
```

INDEXOF(): Returns the first occurrence of a character in a string, searching forward from a start index.

Syntax:

INDEXOF(string, 'character', startIndex)

```
grunt> ans = foreach doc generate (id,name),INDEXOF(name,'a',0);  
grunt> dump ans;  
((1001,Malan),3)  
((1002,Jay),1)  
((1003,lalit),1)
```

```
((1004,Mohit),-1)
((1005,Chauhan),2)
((1006,Suraj),3)
```

LCFIRST(): Converts the first character in a string to lower case.

Syntax:

LCFIRST(expression)

```
grunt> ans = foreach doc generate (id,name),LCFIRST(name);
grunt> dump ans;
((1001,Milan),milan)
((1002,Jay),jay)
((1003,lalit),lalit)
((1004,Mohit),mohit)
((1005,Chauhan),chauhan)
((1006,Suraj),suraj)
```

UCFIRST(): Returns a string with the first character converted to upper case.

Syntax:

UCFIRST(expression)

```
grunt> ans = foreach doc generate (id,hosp),UCFIRST(hosp);
grunt> dump ans;
((1001,apollo),Apollo)
((1002,apollo),Apollo)
((1003,manipal),Manipal)
((1004,columbia),Columbia)
((1005,narayana),Narayana)
((1006,manipal),Manipal)
```

UPPER():Returns a string converted to upper case

Syntax:

UPPER(expression)

```
grunt> ans = foreach doc generate (id,name),UPPER(name);
grunt> dump ans;
((1001,Milan),MILAN)
((1002,Jay),JAY)
((1003,lalit),LALIT)
```

```
((1004,Mohit),MOHIT)
((1005,Chauhan),CHAUHAN)
((1006,Suraj),SURA)
```

LOWER(): Converts all characters in a string to lower case.

Syntax:

LOWER(expression)

```
grunt> ans = foreach doc generate (id,name),LOWER(name);
grunt> dump ans;
((1001,Milan),milan)
((1002,Jay),jay)
((1003,lalit),lalit)
((1004,Mohit),mohit)
((1005,Chauhan),chauhan)
((1006,Suraj),suraj)
```

REPLACE(): To replace existing characters in a string with new characters.

Syntax:

REPLACE(string, 'oldChar', 'newChar');

```
grunt> ans = foreach doc generate (id,hosp),REPLACE(hosp,'apollo','appo');
grunt> dump ans;
((1001,apollo),appo)
((1002,apollo),appo)
((1003,manipal),manipal)
((1004,columbia),columbia)
((1005,narayana),narayana)
((1006,manipal),manipal)
```

BUILT_IN MATH FUNCTIONS

```
$gedit math1.txt
```

```
5
```

```
16
```

```
9
```

```
2.5
```

```
2
```

```
3.5
```

```
3.14
```

```
-2.2
```

```
grunt> mat = load '/home/cloudera/math1.txt' using PigStorage(',') as
                                         (data:float);
```

ABS(): ABSOLUTE VALUE

To get the absolute value of an expression

```
grunt> ans = foreach mat generate data,ABS(data);  
grunt> dump ans;  
(5.0,5.0)  
(16.0,16.0)  
(9.0,9.0)  
(2.5,2.5)  
(2.0,2.0)  
(3.5,3.5)  
(3.14,3.14)  
(-2.2,2.2)
```

CBRT() : cube root

This function is used to get the cube root of an expression.

```
grunt> ans = foreach mat generate data,CBRT(data);  
grunt> dump ans;  
(5.0,1.709975946676697)  
(16.0,2.5198420997897464)  
(9.0,2.080083823051904)  
(2.5,1.3572088082974532)  
(2.0,1.2599210498948732)  
(3.5,1.5182944859378313)  
(3.14,1.464344366810533)  
(-2.2,-1.300591456247907)
```

SQRT() : square root

To get the positive square root of an expression.

```
grunt> ans = foreach mat generate data,SQRT(data);  
grunt> dump ans;  
(5.0,2.23606797749979)  
(16.0,4.0)  
(9.0,3.0)  
(2.5,1.5811388300841898)  
(2.0,1.4142135623730951)  
(3.5,1.8708286933869707)  
(3.14,1.7720045442673602)  
(-2.2,NaN)
```

COS():

This function is used to get the trigonometric cosine of an expression.

```
grunt> ans = foreach mat generate data,COS(data);
grunt> dump ans;
(5.0,0.28366218546322625)
(16.0,-0.9576594803233847)
(9.0,-0.9111302618846769)
(2.5,-0.8011436155469337)
(2.0,-0.4161468365471424)
(3.5,-0.9364566872907963)
(3.14,-0.99999873189461)
(-2.2,-0.5885011558074578)
```

SIN():

To get the sine of an expression.

```
grunt> ans = foreach mat generate data,SIN(data);
grunt> dump ans;
(5.0,-0.9589242746631385)
(16.0,-0.2879033166650653)
(9.0,0.4121184852417566)
(2.5,0.5984721441039564)
(2.0,0.9092974268256817)
(3.5,-0.35078322768961984)
(3.14,0.0015925480124451862)
(-2.2,-0.8084963757576692)
```

TAN():

To get the trigonometric tangent of an angle.

```
grunt> ans = foreach mat generate data,TAN(data);
grunt> dump ans;
(5.0,-3.380515006246586)
(16.0,0.3006322420239034)
(9.0,-0.45231565944180985)
(2.5,-0.7470222972386603)
(2.0,-2.185039863261519)
(3.5,0.3745856401585947)
(3.14,-0.0015925500319664656)
(-2.2,1.37382291908733)
```

CEIL():

This function is used to get the value of an expression rounded up to the nearest integer.

```
grunt> ans = foreach mat generate data,CEIL(data);
grunt> dump ans;
```

(5.0,5.0)
(16.0,16.0)
(9.0,9.0)
(2.5,3.0)
(2.0,2.0)
(3.5,4.0)
(3.14,4.0)
(-2.2,-2.0)

FLOOR():

To get the value of an expression rounded down to the nearest integer.

```
grunt> ans = foreach mat generate data,FLOOR(data);  
grunt> dump ans;
```

(5.0,5.0)
(16.0,16.0)
(9.0,9.0)
(2.5,2.0)
(2.0,2.0)
(3.5,3.0)
(3.14,3.0)
(-2.2,-3.0)

ROUND():

To get the value of an expression rounded to an integer (if the result type is float) or rounded to a long (if the result type is double).

```
grunt> ans = foreach mat generate data,ROUND(data);  
grunt> dump ans;
```

(5.0,5)
(16.0,16)
(9.0,9)
(2.5,3)
(2.0,2)
(3.5,4)
(3.14,3)
(-2.2,-2)

EXP():

This function is used to get the Euler's number e raised to the power of x.

```
grunt> ans = foreach mat generate data, EXP(data);  
grunt> dump ans;
```

(5.0,148.4131591025766)

```
(16.0,8886110.520507872)
(9.0,8103.083927575384)
(2.5,12.182493960703473)
(2.0,7.38905609893065)
(3.5,33.11545195869231)
(3.14,23.103869282414397)
(-2.2,0.1108031530788277)
```

LOG10():

To get the base 10 logarithm of an expression.

```
grunt> ans = foreach mat generate data,LOG10(data);
grunt> dump ans;
(5.0,0.6989700043360189)
(16.0,1.2041199826559248)
(9.0,0.9542425094393249)
(2.5,0.3979400086720376)
(2.0,0.3010299956639812)
(3.5,0.5440680443502757)
(3.14,0.4969296625825472)
(-2.2,NaN)
```

LOG():

To get the natural logarithm (base e) of an expression.

```
grunt> ans = foreach mat generate data,LOG(data);
grunt> dump ans;
(5.0,1.6094379124341003)
(16.0,2.772588722239781)
(9.0,2.1972245773362196)
(2.5,0.9162907318741551)
(2.0,0.6931471805599453)
(3.5,1.252762968495368)
(3.14,1.1442228333291342)
(-2.2,NaN)
```