

Module – 4 (HIVE)

What is Hive?

Hive is a data warehouse infrastructure tool to process structured data in Hadoop. It resides on top of Hadoop to summarize Big Data, and makes querying and analyzing easy.

Initially Hive was developed by Facebook, later the Apache Software Foundation took it up and developed it further as an open source under the name Apache Hive.

Hive primitive data types: tinyint, smallint, int, bigint, string etc.

To switch off safe mode

```
Sudo -u hdfs hdfs dfsadmin -safemode leave
```

To create a data file

```
[cloudera@localhost ~]$ gedit data1;
```

```
1|ramu|20
```

```
2|krish|19
```

```
3|sita|20
```

```
4|john|21
```

```
5|syed|22
```

Save and Close

To connect to hive

```
[cloudera@localhost ~]$ hive
```

To create database

```
hive> create database db2;
```

(to drop database, drop database db2)

To see the list of databases

```
hive> show databases;
```

To make use of the database

```
hive> use db2;
```

To create table

```
hive> create table emp(id int,name string,age int)row format delimited fields  
terminated by '|' lines terminated by '\n' stored as textfile;
```

To see the list of tables

```
hive> show tables;
```

To see the structure of a table

```
hive> describe emp;
```

To load data into hive table from local system

```
hive> load data local inpath '/home/cloudera/data1' into table emp;
```

To see the content of the table

```
hive> select * from emp;
```

To rename the table

```
alter table emp rename to employee;
```

To check the table names

```
hive> show tables;
```

To add new columns to the table

```
hive> alter table employee add columns(city string,pincode int);
```

To see the structure of a table

```
hive> describe employee;
```

To rename the column in a table

```
hive> alter table employee change city addr string;
```

To see the structure of a table

```
hive> describe employee;
```

To select the rows based on condition

```
hive> select * from employee where age>19;
```

To create a table from another table

```
hive> create table empd as select * from employee;
```

```
hive> select * from empd;
```

To limit the number of rows to be displayed

```
hive> select * from employee limit 2;
```

To use built-in functions

```
hive> select upper(name) from employee;
```

```
hive> select count(id) from employee;
```

```
hive> select substr(name,1,3) from employee;
```

To store the output of analysis to some other table

```
hive> insert overwrite table empd select * from employee where age>20;
```

```
hive> select * from empd;
```

To join two tables

```
hive> select * from employee e join empd d on (e.id=d.id);
```

```
hive> select * from employee e left outer join empd d on (e.id=d.id);
```

```
hive> select * from employee e right outer join empd d on (e.id=d.id);
```

To create a view

```
hive> create view emp_v as select id,name from employee;
```

```
hive> select * from emp_v;
```

To store the output of analysis to HDFS file system

```
hive> insert overwrite directory '/user/cloudera/output2'
```

```
select * from employee where age>20;
```

To store the output of analysis to Local file system

```
hive> insert overwrite local directory '/home/cloudera/output1'
```

```
select * from employee where age>20;
```

To quit from hive

```
hive> quit;
```

To check the output file in local system

```
[cloudera@localhost ~]$ ls output1;
```

```
[cloudera@localhost ~]$ cat output1/000000_0
```

To check the output file in Hadoop

```
[cloudera@localhost ~]$ hadoop fs -ls /user/cloudera/
```

```
[cloudera@localhost ~]$ hadoop fs -cat /user/cloudera/output1/000000_0
```

(or)

Through Browser

MODULE 4

HIVE

TO SWITCH OFF SAFE MODE

```
$ sudo -u hdfs hdfs dfsadmin -safemode leave
```

EXP 20: DDL Commands

To Create Two Files On Local Filesystem And Copy It To Hdfs Any Folder

```
[cloudera@localhost ~]$ gedit emp.txt
```

```
1001 | hari | d1 | chennai | 1986-12-10
```

```
1002 | teja | d1 | hyd | 1987-01-21
```

```
1003 | ram | d3 | delhi | 1986-02-11
```

```
1004 | milind | d4 | bang | 1988-03-21
```

```
1005 | jay | d2 | bang | 1988-03-22
```

```
1006 | naveen | d4 | hyd | 1986-04-12
```

```
1007 | naser | d1 | hyd | 1989-11-15
```

```
1008 | rahul | d3 | delhi | 1990-12-23
```

```
[cloudera@localhost ~]$ gedit d.txt
```

```
d1 | research | A-block
```

```
d2 | sales | A-block
```

```
d3 | testing | B-block
```

```
d4 | development | C-block
```

```
[cloudera@localhost ~]$ hadoop fs -put emp.txt /user/cloudera/batch3
```

```
[cloudera@localhost ~]$ hadoop fs -put d.txt /user/cloudera/batch3
```

CONNECT TO HIVE

```
[cloudera@localhost ~]$ hive
```

```
hive>
```

TO CREATE DATABASE

```
hive> ok; (OR)
hive> create database if not exists test;
```

TO LIST OUT DATABASES

```
hive> show databases;
```

TO DROP DATABASE

```
hive> drop database test; (OR)
hive> drop database if exists test; (OR)
hive> drop database if exists test cascade;
```

NOTE!!!! [if exists] & [if not exists] doesn't show error if database already exists while creating time and database doesn't exists while dropping the same.
Without these options , errors displayed clearly.

TO MAKE USE OF THE DATABASE

```
hive> use test;
```

Create Table Statement

Create Table is a statement used to create a table in Hive. The syntax and example are as follows:

Syntax

```
CREATE [TEMPORARY] [EXTERNAL] TABLE [IF NOT EXISTS] [db_name.]
table_name
[(col_name data_type [COMMENT col_comment], ...)]
[COMMENT table_comment]
[ROW FORMAT row_format]
[STORED AS file_format]
```

TO CREATE TABLE

```
hive> create table emp(id int,name string,dept string,place string,dob string)
>comment 'this is employee table'
> row format delimited fields terminated by '|' lines terminated by '\n'
>stored as textfile;
```

(OR) Type IN Single Line

```
hive> create table emp(id int,name string,dept string,place string,dob string)
comment 'this is employee table' row format delimited fields terminated by '|'
lines terminated by '\n' stored as textfile;
```

hive> create table department(did string,dname string,block string) comment 'this is department table' row format delimited fields terminated by '|' lines terminated by '\n' stored as textfile;

NOTE!!! You can mention just mention

Hive>USE test;

Hive> CREATE TABLE emp (.....)

(OR)

Hive> CREATE TABLE test.emp (.....)

TO SEE THE LIST OF TABLES

hive> show tables;

TO SEE THE STRUCTURE OF A TABLE

hive> describe emp;

TO SEE THE STRUCTURE & METADATA INFORMATION OF TABLE

hive> describe formatted emp;

hive> show create table emp;

Contents of directory /user/hive/warehouse/test.db								
Goto : <input type="text" value="/user/hive/warehouse/test.c"/> <input type="button" value="go"/>								
Go to parent directory								
Name	Type	Size	Replication	Block Size	Modification Time	Permission	Owner	Group
department	dir				2018-04-09 23:38	rw-rw-rw-	cloudera	hive
emp	dir				2018-04-09 23:57	rw-rw-rw-	cloudera	hive

Alter Table Statement

It is used to alter a table in Hive.

Syntax

The statement takes any of the following syntaxes based on what attributes we wish to modify in a table.

ALTER TABLE name RENAME TO new_name

ALTER TABLE name ADD COLUMNS (col_spec[, col_spec ...])

ALTER TABLE name CHANGE column_name new_name new_type

ALTER TABLE name REPLACE COLUMNS (col_spec [, col_spec ...])

TO RENAME TABLE NAME

hive> alter table department rename to d;

hive> show tables;

d

emp

TO ADD ONE OR MORE COLUMNS TO THE TABLE


```
hive> alter table d add columns (estb_year int,rating smallint);
```

```
hive> describe d;
```

```
did    string
dname   string
block   string
estb_year  int
rating  smallint
```

TO CHANGE COLUMN NAME OR ITS DATATYPE OR BOTH

```
hive> alter table d change rating rate string;
```

```
hive> describe d;
```

```
did    string
dname   string
block   string
estb_year  int
rate    string
```

```
hive> alter table d change rate rate bigint;
```

```
hive> describe d;
```

```
did    string
dname   string
block   string
estb_year  int
rate    bigint
```

TO REPLACE COLUMNS

```
hive> alter table d replace columns (did string,dname string, block string);
```

```
hive> describe d;
```

```
did    string
dname   string
block   string
```

```
hive> alter table d replace columns (block string);
```

```
hive> describe d;
```

```
block   string
```

```
hive> select * from d;
```

```
d1
d2
d3
d4
```

//if you do REPLACE again, you will get the columns again I,e replace is not removing columns permanently

```
hive> alter table d replace columns (did string,dname string, block string);
```

```
hive> desc d;
```

```
did    string
dname  string
block  string
```

```
hive> select * from d;
```

```
d1    research    A-block
d2    sales    A-block
d3    testing     B-block
d4    development C-block
d5    hr         A-block
```

TO DROP THE TABLE

```
hive> drop table if exists d;          (OR)
```

```
hive> drop table d;
```

EXP 21: Load, Insert of data

Load Data Statement

Generally, after creating a table in SQL, we can insert data using the Insert statement. But in Hive, we can insert data using the LOAD DATA statement.

While inserting data into Hive, it is better to use LOAD DATA to store bulk records. There are two ways to load data: one is from local file system and second is from Hadoop file system.

Syntax

The syntax for load data is as follows:

```
LOAD DATA [LOCAL] INPATH 'filepath' [OVERWRITE] INTO TABLE
tablename
[PARTITION (partcol1=val1, partcol2=val2 ...)]
```

- LOCAL is identifier to specify the local path. It is optional.
- OVERWRITE is optional to overwrite the data in the table.
- PARTITION is optional used table is created with partitions.

TO LOAD FROM LOCAL FILESYSTEM

```
hive> load data local inpath '/home/cloudera/emp.txt' into table emp;
```

```
hive> select * from emp;
```

```
1001 hari    d1    chennai    1986-12-10
1002 teja    d1    hyd        1987-01-21
1003 ram     d3    delhi     1986-02-11
1004 milind  d4    bang     1988-03-21
1005 jay     d2    bang     1988-03-22
1006 naveen  d4    hyd      1986-04-12
1007 naser   d1    hyd      1989-11-15
1008 rahul   d3    delhi    1990-12-23
```

Contents of directory /user/hive/warehouse/test.db/emp								
Goto : <input type="text" value="/user/hive/warehouse/test.c"/> <input type="button" value="go"/>								
Go to parent directory								
Name	Type	Size	Replication	Block Size	Modification Time	Permission	Owner	Group
emp.txt	file	238 B	3	128 MB	2018-04-09 23:57	rw-rw-rw-	cloudera	supergroup

TO LOAD FROM HADOOP FILE SYSTEM

```
hive> load data inpath '/user/cloudera/emp.txt' into table emp;
```

```
hive> select * from emp;
```

```
1001 hari d1 chennai 1986-12-10
1002 teja d1 hyd 1987-01-21
1003 ram d3 delhi 1986-02-11
1004 milind d4 bang 1988-03-21
1005 jay d2 bang 1988-03-22
1006 naveen d4 hyd 1986-04-12
1007 naser d1 hyd 1989-11-15
1008 rahul d3 delhi 1990-12-23
1001 hari d1 chennai 1986-12-10
1002 teja d1 hyd 1987-01-21
1003 ram d3 delhi 1986-02-11
1004 milind d4 bang 1988-03-21
1005 jay d2 bang 1988-03-22
1006 naveen d4 hyd 1986-04-12
1007 naser d1 hyd 1989-11-15
1008 rahul d3 delhi 1990-12-23
```

Contents of directory /user/hive/warehouse/test.db/emp								
Goto : <input type="text" value="/user/hive/warehouse/test.c"/> <input type="button" value="go"/>								
Go to parent directory								
Name	Type	Size	Replication	Block Size	Modification Time	Permission	Owner	Group
emp.txt	file	238 B	3	128 MB	2018-04-09 23:48	rw-rw-rw-	cloudera	supergroup
emp_copy_1.txt	file	238 B	3	128 MB	2018-04-09 22:25	rw-rw-rw-	cloudera	cloudera

NOTE!!!!Here, Since no overwrite was used; the data got appended to same table. And in Hive/warehouse/test.db two copies of same content got generated.

TO LOAD USING OVERWRITE KEYWORD

```
hive> load data local inpath '/home/cloudera/emp.txt' overwrite into table emp;
```

```
hive> select * from emp;
```

```
1001 hari d1 chennai NULL
1002 teja d1 hyd NULL
```

```

1003 ram d3 delhi NULL
1004 milind d4 bang NULL
1005 jay d2 bang NULL
1006 naveen d4 hyd NULL
1007 naser d1 hyd NULL
1008 rahul d3 delhi NULL

```

Contents of directory /user/hive/warehouse/test.db/emp								
Goto : <input type="text" value="/user/hive/warehouse/test.c"/> <input type="button" value="go"/>								
Go to parent directory								
Name	Type	Size	Replication	Block Size	Modification Time	Permission	Owner	Group
emp.txt	file	238 B	3	128 MB	2018-04-10 00:19	rw-rw-rw-	cloudera	supergroup

NOTE!!! Once the data is loaded to hive table from hadoop filesystem, the file "emp.txt" no more exists in /user/cloudera I.e loading from hadoop filesystem is like cut and paste to hive; whereas its like copy & paste when loaded from local filesystem.

So , if you have loaded from hadoop filesystem once, then you can't load or load with overwrite to hive table from hadoop filesystem again.... Because you will get ERROR: "invalid path as file is cut already from that location."

```

hive> load data local inpath '/home/cloudera/d.txt' overwrite into table department;
hive> select * from department;
d1    research    A-block
d2    sales       A-block
d3    testing     B-block
d4    development C-block

```

Contents of directory /user/hive/warehouse/test.db/department								
Goto : <input type="text" value="/user/hive/warehouse/test.c"/> <input type="button" value="go"/>								
Go to parent directory								
Name	Type	Size	Replication	Block Size	Modification Time	Permission	Owner	Group
d.txt	file	79 B	3	128 MB	2018-04-10 03:00	rw-rw-rw-	cloudera	supergroup

EXP 22: BUILT_IN RELATIONAL OPERATORS

A=B, A<B, A<=B, A>B, A>=B, A IS NULL, A IS NOT NULL, A LIKE B, A!=B

BUILT IN ARTHMETIC OPERATORS

A+B, A-B, A*B, A/B, A%B, A&B, A|B, A^B, ~A
A&B : bitwise and operation
A|B: bitwise or operation
A^B;bitwise XOR operation
~A: bitwise not operation

BUILT IN LOGICAL OPERATORS

A AND B, A OR B, NOT A, A || B, A&&B, !A

Q)Display details of employee whose employee id is greater than and equal to 1003 and doesn't come from hyd.

hive> select * from emp where id >=1003 and place !='hyd'; (OR)

hive> select * from emp where id >=1003 and place not in ('hyd');

1003	ram	d3	delhi	1986-02-11
1004	milind	d4	bang	1988-03-21
1005	jay	d2	bang	1988-03-22
1008	rahul	d3	delhi	1990-12-23

Q)Display details of department whose id is less than d2 or department name is development.

hive> select * from department where did<'d2' or dname='development';

(OR)

hive> select * from department where did<'d2' or dname like 'development';

(OR)

hive> select * from department where did<'d2' or dname like 'd%';

d1	research	A-block
d4	development	C-block

Q)Display details of department whose department name's second letter is 'e'.

hive> select * from department where did<'d2' or dname like '_e%';

d1	research	A-block
d3	testing	B-block
d4	development	C-block

Q)Display total no:of employees ,minimum of their employee no, avg of their employee no,max of their employee, sum of their employee from employee dataset.

hive> select count(*),min(id),avg(id),max(id),sum(id) from emp;

8	1001	1004.5	1008	8036
---	------	--------	------	------

Q) Display department id, count of employees in each department

hive> select count(*),dept from emp group by dept;

3	d1
1	d2
2	d3
2	d4

Q) Display department id,count of employees in each department and display rows those have count more than 2.

```
hive> select count(*),dept from emp group by dept having count(*)>2;
3      d1
```

Q) Display department id,count of employees in each department in descending order of count.

```
hive> select count(*) as c,dept from emp group by dept order by c desc;
3      d1
2      d3
2      d4
1      d2
```

Q) Display department id,count of employees in each department in descending order of count and display only first two rows.

```
hive> select count(*) as c,dept from emp group by dept order by c desc limit 2;
3      d1
2      d3
```

EXP 23: TO JOIN TWO TABLES

```
hive> select * from emp e join department d on (e.dept=d.did);
001  hari  d1  chennai  1986-12-10  d1  research  A-block
1002  teja  d1  hyd  1987-01-21  d1  research  A-block
1007  naser d1  hyd  1989-11-15  d1  research  A-block
1005  jay   d2  bang  1988-03-22  d2  sales  A-block
1003  ram   d3  delhi 1986-02-11  d3  testing B-block
1008  rahul d3  delhi 1990-12-23  d3  testing B-block
1004  milind d4  bang  1988-03-21  d4  development C-block
1006  naveen d4  hyd  1986-04-12  d4  development C-block
hive> select * from emp e left outer join department d on (e.dept=d.did);
1001  hari  d1  chennai  1986-12-10  d1  research  A-block
1002  teja  d1  hyd  1987-01-21  d1  research  A-block
1007  naser d1  hyd  1989-11-15  d1  research  A-block
1005  jay   d2  bang  1988-03-22  d2  sales  A-block
1003  ram   d3  delhi 1986-02-11  d3  testing B-block
1008  rahul d3  delhi 1990-12-23  d3  testing B-block
1004  milind d4  bang  1988-03-21  d4  development C-block
1006  naveen d4  hyd  1986-04-12  d4  development C-block
1009  jay   d6  hyd  1988-07-19  null  null  null
```

NOTE!!if d6 department not there ,then no matching on right side table values

```
hive> select * from emp e right outer join department d on (e.dept=d.did);
1001 hari d1 chennai 1986-12-10 d1 research A-block
1002 teja d1 hyd 1987-01-21 d1 research A-block
1007 naser d1 hyd 1989-11-15 d1 research A-block
1005 jay d2 bang 1988-03-22 d2 sales A-block
1003 ram d3 delhi 1986-02-11 d3 testingB-block
1008 rahul d3 delhi 1990-12-23 d3 testingB-block
1004 milind d4 bang 1988-03-21 d4 development C-block
1006 naveen d4 hyd 1986-04-12 d4 development C-block
NULL NULL NULL NULL NULL d5 hr A-block
```

EXP 24: TO CREATE A VIEW

```
hive> create view emp_v as select id,name from emp where id>1003;
```

```
hive> select * from emp_v;
```

```
1004 milind
1005 jay
1006 naveen
1007 naser
1008 rahul
```

TO DROP THE VIEW

```
hive> drop view emp_v;
```

To Use Built-In Functions

```
hive> select upper(name) from emp;
```

```
HARI
TEJA
RAM
MILIND
JAY
NAVEEN
NASER
RAHUL
```

```
hive> select count(id) from emp;
```

```
8
```

```
hive> select substr(name,1,3) from emp;
```

```
har
tej
ram
mil
jay
nav
nas
```

rah

hive> select substr(name,2) from emp;

*ari
eja
am
ilind
ay
aveen
aser
ahul*

hive> select substr(name,3,2) from emp;

*ri
ja
m
li
y
ve
se
hu*

Syntax: substr(string,starting index,no of character)

Note: if no of characters not mentioned then it returns from the start position to the end of the string

TO CREATE A TABLE FROM ANOTHER TABLE

**hive> create table abc as select * from emp;
hive> select * from abc;**

TO STORE THE OUTPUT OF ANALYSIS TO SOME OTHER TABLE

**hive> insert overwrite table abc select * from emp where id>1003;
hive> select * from abc;**

004	Milind	D4	Bang	1988-03-21
1005	Jay	D2	Bang	1988-03-22
1006	Naveen	D4	Hyd	1986-04-12
1007	Naser	D1	Hyd	1989-11-15
1008	Rahul	D3	Delhi	1990-12-23

Note:schema should match(No.of column should match)

To Store The Output Of Analysis To Hdfs File System

**Hive> Insert Overwrite Directory '/User/Cloudera/Output1' Select * From Emp
Where Id>1003;**

Note!! Where Output1 Is A New Directory In /User/Cloudera, Which Will Get Created Automatically

To Check The Output File In Hadoop File System

[Cloudera@Localhost ~]\$ Hadoop Fs -Ls /User/Cloudera/Output1

[Cloudera@Localhost ~]\$ Hadoop Fs -Cat /User/Cloudera/Output1/000000_0
(Or)

Hive> Dfs -Ls /User/Cloudera/Output1;

-Rw-R--R-- 3 Cloudera Supergroup 149 2018-04-17 03:32
/User/Cloudera/Output1/000000_0

Hive> Dfs -Cat /User/Cloudera/Output1/0*;

1004milindd4bang1988-03-21

1005jayd2bang1988-03-22


1006naveend4hyd1986-04-12

1007naserd1hyd1989-11-15

1008rahuld3delhi1990-12-23

(Or)

Use Browser Of Your Vm To Go To The Location And Verify The File



TO STORE THE OUTPUT OF ANALYSIS TO LOCAL FILE SYSTEM

hive> insert overwrite local directory '/home/cloudera/output1' select * from employee where id>1003;

TO CHECK THE OUTPUT FILE IN LOCAL SYSTEM

- [cloudera@localhost ~]\$ ls
 - [cloudera@localhost ~]\$ cd output1
 - [cloudera@localhost ~]\$ cat 000000_0 (OR) \$gedit 000000_0
- (OR)
- [cloudera@localhost ~]\$ ls output1
 - [cloudera@localhost ~]\$ cat output1/000000_0



```
000000_0 (~/.output1) - gedit
File Edit View Search Tools Documents Help
Open Save Print Undo Redo Cut Copy Paste
000000_0
1004;milind;d4;bang;1988-03-21
1005;jay;d2;bang;1988-03-22
1006;naveen;d4;hyd;1986-04-12
1007;naser;d1;hyd;1989-11-15
1008;rahul;d3;delhi;1990-12-23
```

TO QUIT FROM HIVE

hive> quit;